

Job Intelligence Report

Target Company: EY GDS -- Staff GenAI Engineer (Pune, 3+ yrs)

Generated: May 01, 2026 at 07:41 AM

Section	Content
1	JD Analysis — Fit Score, Keywords, Gaps
2	Company Intelligence — Culture, Salary, Process
3	Interview Preparation Kit

SECTION 1 — JD ANALYSIS

FIT SCORE

78/100 -- Strong technical match on mandatory skills: LangChain + LangGraph, FastAPI, RAG, vector databases, prompt engineering, AWS, Python OOP. Minor gaps: DynamoDB not in CV, JWT auth not mentioned, AWS Bedrock not hands-on. Pranav is slightly overlevelled (5+ years vs. 3+ years required) which means he'd enter at Staff not Senior -- manageable. EY GDS is a delivery-focused environment, and the role emphasises implementation within defined architecture rather than design ownership.

TOP JD KEYWORDS

- RAG pipelines (based on provided architecture)
- Prompt engineering and iteration
- FastAPI and REST API integration
- LangChain and LangGraph
- Vector databases and embeddings
- AWS Bedrock and Azure OpenAI
- JWT authentication (Python)
- AWS deployment and CI/CD standards
- Python OOP (strong)
- LLMOps / evaluation and monitoring (good to have)

MATCHING SKILLS

- LangChain + LangGraph: Both explicitly on CV; LangGraph used in 6-agent production pipeline -- exceeds the "hands-on experience" requirement
- FastAPI: Core framework used across all projects; 15+ endpoints including async background tasks and SSE streaming -- strong match
- RAG: Production RAG with 5-tier extraction, cross-encoder reranking, HyDE, multilingual support, RAGAS evaluation -- far exceeds "implement RAG based on provided architecture"
- Vector databases + embeddings: Pinecone, ChromaDB, FAISS, BAAI/bge-m3 embeddings -- direct match
- Python OOP (strong): Advanced Python throughout; Pydantic, asyncio, FastAPI -- strong OOP patterns
- Prompt engineering: System prompts, chain-of-thought, difficulty rubrics, decomposition -- matches "prompt iteration under defined standards"
- AWS: EC2, S3 working knowledge; Azure AI Fundamentals certified
- API integration: FastAPI REST APIs integrated into enterprise document and chatbot pipelines
- LLMOps / evaluation: RAGAS, LLM-as-judge, faithfulness scoring, logging via L1-L4 quality indicators -- matches "good to have" evaluation and monitoring
- Enterprise chatbots: RAG-based enterprise chatbot delivered at Optimum -- direct match

SKILL GAPS

- AWS DynamoDB: Not in CV -- PostgreSQL and S3 present; DynamoDB is AWS's NoSQL; easy to learn but not hands-on yet
- JWT authentication: Not mentioned in CV -- standard OAuth2/JWT implementation; acknowledge gap and confirm it's a quick learn given FastAPI experience
- AWS Bedrock: Not hands-on -- used OpenAI and Anthropic APIs directly, not via Bedrock; Azure AI Fundamentals certified shows cloud ML awareness
- Azure OpenAI: Not explicitly hands-on -- used Anthropic Claude and OpenAI APIs directly; Azure OpenAI is same models via Azure endpoint, minimal learning curve
- 3+ year role: Pranav is 5+ years -- slightly overqualified; will come in at Staff, not Senior or Principal. Position as "I want to apply my depth to high-quality delivery in a structured enterprise environment"

POINTS TO EMPHASISE

- "I already have LangChain AND LangGraph in production -- your mandatory skills are cleared before the interview starts"
- FastAPI is your strongest deployment framework -- 15+ endpoints, async tasks, SSE streaming -- well beyond what this delivery role requires
- The "good to have" LLMOps section (RAGAS, monitoring, evaluation) is something you've already built -- position this as exceeding expectations
- EY's structured delivery model actually suits someone with your project delivery track record at Optimum
- RAG depth (5-tier extraction, reranking, RAGAS) is well above the "implement RAG per architecture" requirement -- shows you can improve the architecture, not just follow it

ROLE SUMMARY

EY GDS (Global Delivery Services) is EY's offshore delivery centre for the Big 4 firm's global consulting practice. The Staff GenAI Engineer role is an implementation-focused position that builds GenAI components to architectures designed by senior engineers and managers. This is NOT a design-ownership role -- it's a quality execution role. That framing matters: EY is looking for engineers who can implement reliably within standards, communicate well, and iterate on prompts and pipelines with discipline. For Pranav, this is a slightly underlevelled role technically, but the EY Big 4 brand on a CV has real career value. Pune GDS is a well-established delivery centre.

SECTION 2 — COMPANY INTELLIGENCE

COMPANY OVERVIEW

EY (Ernst & Young) is one of the Big 4 professional services firms, with revenue exceeding \$49B globally (FY2024) and 395,000+ employees across 150+ countries. EY GDS (Global Delivery Services) is EY's offshore delivery network, with major centres in Pune, Bengaluru, Chennai, Hyderabad, and Kochi. The Pune GDS centre (located in Hinjewadi) is one of the largest, serving EY's global consulting, assurance, and technology practices. The AI & Analytics (AIA) team within EY GDS is the unit this GenAI role sits in -- building AI components for EY's global client engagements.

CULTURE & WORK ENVIRONMENT

- Glassdoor India (EY GDS): 3.8/5 overall; 3.6/5 work-life balance; 4.0/5 culture & values
- 75% of India employees would recommend EY to a friend
- Large, structured corporate environment with defined career tracks (Staff → Senior → Manager → Senior Manager → Director)
- Hybrid work model: typically 3 days in office in Hinjewadi, Pune; some flexibility
- EY has strong internal learning: EY Badges, EY Tech MBA, LinkedIn Learning -- active learning culture
- Work-life balance varies by project: during EY's busy season (Jan-April) hours can extend significantly
- GenAI/AIA team is newer and faster-moving than traditional assurance teams
- Reviews note good benefits package (health insurance, meal cards, cab facility) and professional work environment
- Career progression is structured -- clear criteria for promotion from Staff to Senior

SALARY RANGES (EY GDS Pune, Staff GenAI Engineer, 3-5 years)

- Staff Engineer, GenAI (3-5 years): ■18-26 LPA
- Senior Staff / Senior Engineer (5-7 years): ■24-32 LPA
- Pranav enters at Staff given the 3+ years JD, but with 5+ years could negotiate Senior Staff placement: ■22-28 LPA target
- EY GDS is slightly above mid-market for Indian delivery centres; below product companies by 25-40%
- Benefits package is comprehensive: medical, gratuity, PF, meal allowance, cab facility
- Variable component: 15-20% of CTC at Staff level; performance-linked

INTERVIEW PROCESS

EY GDS AIA team GenAI hiring process:

1. HR screening: background, EY values alignment, salary, notice period (30 min)

2. Technical Round 1: Python OOP, FastAPI, GenAI fundamentals (RAG, chunking, CoT, guardrails) -- 45 min
 3. Technical Round 2: LangChain/LangGraph implementation, system design, past project walkthrough -- 60 min
 4. Manager interview: delivery mindset, communication, ability to work within defined standards -- 30 min
 5. HR offer discussion
- EY values STAR-format answers in behavioural rounds -- prepare structured responses
 - Process typically takes 3-4 weeks for GDS hiring
 - Technical rounds may include live coding in Python (FastAPI endpoint, RAG component implementation)

PROS

- Big 4 brand -- EY on a CV opens doors globally; valuable for long-term career mobility
- Structured career track -- Staff → Senior → Manager is a clear progression with defined criteria
- Strong learning culture: EY Badges, Tech MBA equivalent, certification funding
- Exposure to global Fortune 500 clients' AI needs through EY's consulting practice
- Good benefits: comprehensive health insurance, PF, meal allowance, cab from Pune
- Pune GDS is large and stable -- not a startup risk
- GenAI/AIA team is growing -- good timing to join an expanding practice

CONS / RED FLAGS

- Staff-level role means implementation work, not design ownership -- could feel underutilised given Pranav's depth
- Delivery centre model: pace set by global senior team; India team executes
- Big 4 process-heavy culture -- slower than startups, more approvals, more documentation
- Salary below product companies and some IT services firms for equivalent technical depth
- Busy season (Jan-April) means long hours if supporting global assurance practices
- Hinjewadi office location: good cab service, but Pune traffic during in-office days

RECENT NEWS & DEVELOPMENTS

- EY launched EY.ai platform (2024) -- centralised AI platform for client and internal use
- Active investment in GenAI practice: hiring 20,000 AI-skilled professionals globally
- EY GDS AIA team expanding rapidly in Pune and Bengaluru through 2025-2026
- Partnership with Microsoft (Azure OpenAI) and AWS (Bedrock) for enterprise AI delivery
- EY's AI business expected to generate \$10B+ revenue by FY2026

TECH STACK (EY GDS AIA team)

-
- LLMs: Azure OpenAI (GPT-4), AWS Bedrock (Claude, Llama), Anthropic Claude
 - Frameworks: LangChain, LangGraph (per this JD), LlamaIndex
 - Vector DBs: Pinecone, FAISS, Azure AI Search, OpenSearch
 - Cloud: AWS (primary for this role), Azure (for some client engagements)
 - Auth: JWT, OAuth2 (via FastAPI)
 - Databases: AWS DynamoDB, PostgreSQL, S3
 - CI/CD: GitHub Actions, AWS CodePipeline
 - Monitoring: CloudWatch, LangSmith (in some engagements), custom logging

NEGOTIATION INTEL

- EY GDS has banded salary structures with limited flex on base -- typically 8-12% negotiation room
- Target ask: ■24-28 LPA -- position for Senior Staff placement given 5+ years vs. 3+ years JD
- Script: "Given my 5+ years of experience and hands-on LangGraph and RAG experience that exceeds the Staff level requirements, I'd like to discuss Senior Staff placement at ■24-28 LPA rather than base Staff band"
- Ask specifically: "What band does this role sit in and what is the criteria for Senior Staff vs. Staff placement?"
- EY has a good learning budget -- if base is capped, ask for certification reimbursement (AWS Solutions Architect, Azure AI Engineer, etc.)
- EY's joining bonus is rare for GDS; focus negotiation on band placement and learning budget instead

SECTION 3 — INTERVIEW PREPARATION KIT

TECHNICAL QUESTIONS (likely to be asked)

1. Write a FastAPI endpoint that accepts a document, runs it through a RAG pipeline, and returns a streamed response. -- Tests: Core FastAPI + RAG implementation; answer with your SSE streaming endpoint architecture
2. How do you implement chunking for a RAG pipeline? What are the tradeoffs between fixed-size, semantic, and sentence-window chunking? -- Tests: RAG fundamentals; fixed-size = simple, loses context at boundaries; semantic = better coherence, more compute; sentence-window = good for conversational retrieval
3. Walk me through a complete prompt engineering iteration cycle. How do you know when a prompt is good enough? -- Tests: Prompt iteration discipline; answer: define eval criteria first, test on representative samples, measure faithfulness/relevance, iterate on system prompt and few-shots until threshold met
4. What is Chain-of-Thought prompting and when would you use it vs. direct prompting? -- Tests: Prompt techniques; CoT = ask model to reason step-by-step before answering; use for multi-step reasoning, math, complex decisions; not needed for simple classification
5. How do you implement guardrails in a GenAI pipeline? What specifically do you check on input vs. output? -- Tests: Safety implementation; input: PII detection, injection attacks, off-topic queries; output: faithfulness check, PII in response, toxic content filter
6. You're asked to integrate an LLM with a FastAPI service using JWT authentication. Walk me through the approach. -- Tests: Gap area; honest answer: "I've built FastAPI services extensively -- JWT auth is a standard OAuth2 bearer token flow in FastAPI using python-jose and passlib; I can implement it; I haven't done it in a GenAI-specific context but the pattern is standard"
7. What is the difference between AWS Bedrock and calling OpenAI directly? When would you use Bedrock? -- Tests: AWS knowledge; Bedrock = AWS-managed foundation model API (Claude, Llama, Titan); use when client is AWS-native and wants to keep data within AWS; same models, different endpoint and auth
8. How do you add logging and monitoring to a GenAI pipeline? What signals matter? -- Tests: LLMOps good-to-have; answer: log per-request latency, token usage, faithfulness score, retrieval quality, error rate; your L1-L4 quality indicator system is a direct answer
9. EY works in defined architectures -- if you're given a RAG architecture spec to implement, what would you clarify before writing code? -- Tests: Delivery mindset; chunking strategy, embedding model, retrieval top-K, reranking yes/no, evaluation criteria, acceptance thresholds -- clarify upfront to avoid rework
10. How do you do code review for a GenAI component? What would you look for that's different from standard code review? -- Tests: Quality mindset; standard: code style, error handling, tests; GenAI-specific: prompt injection risks, determinism/temperature settings, eval metrics coverage, hallucination edge cases

BEHAVIOURAL QUESTIONS (EY GDS culture-specific)

1. EY values working within defined standards. Tell me about a time you implemented a solution within architectural constraints set by someone else.
2. Describe a time you caught a quality issue in a GenAI component during code review or testing. How did you handle it?
3. EY serves global clients -- tell me about a time you collaborated with remote or cross-functional stakeholders on a delivery.
4. EY has a strong learning culture. What have you learned in the last 6 months that made you better at your job?
5. Tell me about a time you had to document a complex AI solution clearly for a non-technical audience.

STAR STORY SUGGESTIONS

Story 1: "Implementing RAG within delivery standards"

- Situation: Enterprise client needed document QA system; requirements came from client stakeholder with specific accuracy and reliability expectations
- Task: Build and deliver a RAG pipeline that met the client's defined quality criteria
- Action: Implemented 5-tier extraction, embedding, retrieval, reranking, and generation pipeline; set up RAGAS evaluation framework to measure against client's quality threshold; iterated on chunking and reranking parameters until faithfulness > 0.85
- Result: System delivered within timeline; client validated against quality criteria; documented architecture and evaluation results for handover
- Best used for: "Working within defined standards", "Quality mindset", "Delivery discipline"

Story 2: "FastAPI RAG service with production monitoring"

- Situation: Needed to expose a RAG pipeline as a production REST API with quality indicators for enterprise use
- Task: Build a FastAPI service that handled async ingestion, SSE streaming, and per-request quality reporting
- Action: Implemented 15+ FastAPI endpoints with async background tasks, SSE streaming, and L1-L4 quality indicators (faithfulness, relevance, latency, hallucination risk) per request
- Result: Production service containerised with Docker Compose; quality indicators give operators visibility into each response's reliability
- Best used for: FastAPI questions, production thinking, monitoring

Story 3: "Prompt iteration and evaluation at Upthink"

- Situation: Clients had diverse NLP problems requiring different model configurations and prompt strategies
- Task: Deliver reliable AI implementations across 500+ client use cases with different requirements
- Action: Developed systematic approach to prompt iteration: define evaluation criteria first, test on representative samples, iterate on system prompt and examples until threshold met, document results

- Result: Consistent delivery quality across diverse client types; strong reputation for reliable AI implementations
- Best used for: "Prompt engineering discipline", "Consistent delivery", "Communication"

SALARY NEGOTIATION STRATEGY

- Target ask: ■24-28 LPA (position for Senior Staff, not base Staff)
- Walk-away: ■18 LPA
- Script: "This JD is for 3+ years and I have 5+ years -- including production LangGraph and RAG systems that exceed Staff-level complexity. I'd like to discuss Senior Staff placement at ■24-28 LPA, which reflects the experience level I'm bringing."
- EY has banded structure -- ask: "What band does this role sit in and is Senior Staff placement possible given my experience?"
- If base is capped at Staff: negotiate learning budget (AWS certification, Azure AI Engineer) and ask about accelerated promotion timeline
- EY pays good non-cash benefits -- confirm: medical insurance coverage, PF, meal allowance, cab facility

SMART QUESTIONS TO ASK

1. "This role is implementation-focused -- as I grow, what's the path from Staff to Senior where I can take on more architecture and design ownership?"
2. "What AWS services does this team use most? Is the primary stack Bedrock + DynamoDB, or is there a mix?"
3. "What does a typical sprint look like in the AIA team -- what's the review and quality gate process for GenAI components?"
4. "EY recently launched the EY.ai platform -- does this role contribute to that or to specific client engagements?"
5. "What certifications or learning programs are supported for Staff engineers joining the AIA practice?"

KEY TALKING POINTS

1. "I already have LangChain AND LangGraph in production -- your two mandatory framework requirements are cleared before we start discussing; this isn't potential, it's shipped code"
2. "My FastAPI experience goes beyond basics -- async background ingestion, SSE streaming, 15+ endpoints -- I can implement the API layer you need reliably and cleanly"
3. "I've built the 'good to have' LLMOps items (RAGAS evaluation, monitoring, logging) as standard practice in my projects -- I bring evaluation discipline, not just implementation speed"