

Job Intelligence Report

Target Company: NTT DATA -- Mid GenAI Engineer (Pune, 5+ yrs)

Generated: May 01, 2026 at 07:41 AM

Section	Content
1	JD Analysis — Fit Score, Keywords, Gaps
2	Company Intelligence — Culture, Salary, Process
3	Interview Preparation Kit

SECTION 1 — JD ANALYSIS

FIT SCORE

85/100 -- Exceptional match on the Agentic AI and RAG stack. LangGraph, CrewAI, LangChain, Pinecone, hybrid search, re-ranking, guardrails, PII redaction, LLM evaluation -- all present and demonstrated with real shipped projects. The resume selection checklist (2+ GenAI/Agentic solutions, 5+ years Python, 3-5 client POCs) is cleared comfortably. Gaps are Semantic Kernel, AutoGen, LangSmith, and cloud depth beyond AWS working knowledge.

TOP JD KEYWORDS

- Agentic systems (workflow/graph agents, tool use, memory)
- LangGraph, LangChain, CrewAI (explicitly listed)
- Hybrid search + vector stores + re-ranking
- LLM evaluation and observability
- Guardrails, PII redaction, policy filters, agent observability
- Production MLOps (min 2 models deployed)
- Chunking strategies and continuous data quality
- Open-source LLMs (Ollama, HuggingFace)
- Python 5+ years
- Hyperscaler deployment (Azure/AWS/GCP)

MATCHING SKILLS

- LangGraph: 6-agent StateGraph pipeline with supervisor routing, parallel Send API fan-out, and inter-agent feedback loops -- direct, production-level experience matching their #1 framework requirement
- CrewAI: 4-agent Job Intelligence System -- JD Analyst, Company Researcher, CV Tailor, Interview Coach
- LangChain: Listed in skills; used as foundation for LangGraph pipelines
- MCP (Model Context Protocol): 4 custom stdio servers -- differentiator beyond the JD requirements
- Hybrid search + re-ranking: Enterprise Q&A Agent uses cross-encoder reranking (bi-encoder recall + cross-encoder precision) -- exact match
- Vector stores: Pinecone (cloud), ChromaDB (local), FAISS (familiar) -- multi-store experience
- LLM evaluation: RAGAS, LLM-as-judge, faithfulness scoring, hallucination detection -- maps directly to "evaluation & observability" requirement
- Guardrails + PII redaction: Input/output guardrails with PII redaction implemented in Enterprise Q&A Agent -- explicit match
- Agent observability: L1-L4 quality indicators in Streamlit UI -- production monitoring
- Open-source LLMs: LLaMA/Mistral via Ollama -- hands-on open-source model deployment

- 2+ models in production: Dental (95% accuracy), cotton quality (92%), RAG system, LLM evaluator -- well above the minimum 2 threshold
- Python 5+ years: Core language across all work -- clear match
- Client POCs/MVPs: Multiple delivered at Optimum Data Analytics across healthcare, BFSI, industrial
- NLP + CV + DL: Full stack -- sentiment analysis, classification, Vision Transformers, YOLO, CNNs

SKILL GAPS

- Semantic Kernel: Not on CV -- Microsoft's framework, minor gap; CrewAI and LangGraph cover the same orchestration space
- AutoGen: Familiar but not hands-on -- minor gap; CrewAI experience transfers conceptually
- LangSmith: Not explicitly listed -- RAGAS and LLM-as-judge cover the same evaluation purpose; be ready to say "I use RAGAS for evaluation; learning LangSmith is straightforward given my LangGraph depth"
- Cloud depth: AWS EC2/S3 working knowledge, Azure AI Fundamentals; JD wants hyperscaler experience -- not deep IaaS experience, but borderline
- Fine-tuning/FineTuning LLMs: Not hands-on -- FinBERT used (pre-trained domain model), but no PEFT/LoRA fine-tuning; frame FinBERT as "domain-adapted model selection" rather than fine-tuning
- DevOps beyond Docker: Docker Compose is solid; CI/CD pipelines and Kubernetes not mentioned
- LangSmith / human-in-the-loop: RAGAS covers evaluation; human-in-the-loop workflows not explicitly built

POINTS TO EMPHASISE

- Lead with the resume selection checklist NTT DATA published: 2 GenAI/Agentic solutions (Contract Intelligence + Enterprise RAG), 5+ years Python, 3-5 client POCs (Optimum projects) -- tick every box explicitly in your cover note
- MCP is a differentiator beyond what the JD asks: "I've built 4 custom MCP studio servers -- not many engineers in India have done this"
- LangGraph Send API for parallel fan-out shows production-level graph orchestration, not just tutorial knowledge
- RAGAS + LLM-as-judge + PII redaction + guardrails = you've already built what their "guardrails and observability" requirement is asking for
- Open-source model experience (LLaMA/Mistral via Ollama) -- exact match to "working experience with open-source models"
- 95% accuracy medical imaging model + 92% quality detection = production MLOps with quantified outcomes

ROLE SUMMARY

NTT DATA is building enterprise GenAI solutions for global Fortune 100 clients across multiple industries. This mid-level role sits in their India delivery team (Pune) and requires engineers who can independently ship agentic pipelines -- chat copilots, workflow agents, RAG systems -- while managing client stakeholders across

geographies. They want someone who can bridge research and production: understand emerging GenAI frameworks, apply them reliably, and evaluate their own systems. Pranav's profile matches the technical bar, clears the published resume checklist, and brings the rare combination of LangGraph + MCP + production evaluation that makes him stand out in this applicant pool.

SECTION 2 — COMPANY INTELLIGENCE

COMPANY OVERVIEW

NTT DATA is a \$30 billion global IT services and consulting firm, part of the NTT Group which invests over \$3.6 billion annually in R&D. The company serves 75% of the Fortune Global 100, operates in 50+ countries, and employs approximately 190,000 people worldwide. NTT DATA's India operations -- headquartered in Pune and Bangalore -- are a major delivery hub for global clients, with Pune being one of their largest centres. Their AI and Data practice is growing rapidly, with GenAI solutions a core GTM offering.

CULTURE & WORK ENVIRONMENT

- Glassdoor India rating: 3.8/5 overall; 3.6/5 work-life balance
- Culture described as structured and professional -- typical large IT services environment
- Pune office is well-established (10,000+ employees in India total) with strong infrastructure
- Hybrid/in-office model; client-site visits may be required for specific engagements
- Strong internal learning programs: NTT DATA University, certification reimbursement
- Multi-geography client collaboration is core -- expect calls with US/Europe stakeholders
- Slower decision-making than startups; processes are mature but bureaucratic at times
- Employee reviews note stable employment and good benefits but limited fast-track growth

SALARY RANGES (Pune, GenAI Engineer, mid-level, 5+ years)

- ML/AI Engineer 4-6 years (IT services): ■18-26 LPA
- GenAI Engineer with Agentic AI (5+ years, Pune): ■24-34 LPA
- Agentic AI / LangGraph premium in IT services: +15-25% over standard ML roles
- NTT DATA typically pays at the mid-range of IT services bands; not top-of-market but above TCS/Infosys
- Variable component: 15-20% of CTC

INTERVIEW PROCESS

For experienced GenAI hires at NTT DATA (Pune delivery centre):

1. HR screening: background, notice period, salary expectations (30 min)
2. Technical Round 1: Python, ML fundamentals, GenAI concepts, RAG architecture (45-60 min)
3. Technical Round 2: Agentic AI system design, LangGraph/CrewAI depth, live problem discussion (45-60 min)
4. Manager/Client Panel: Cultural fit, client communication ability, past project delivery (30-45 min)
5. HR offer discussion: Compensation, joining date, benefits

- Process typically takes 2-4 weeks
- Walk-in formats sometimes used for bulk hiring drives; may get same-day technical decisions
- NTT DATA values clear communication about multi-geography stakeholder experience

PROS

- \$30B brand -- globally recognised, strong CV value for future moves
- Fortune 100 client exposure -- real enterprise-scale AI problems
- Pune base -- no relocation needed
- Strong internal training and certification budget
- Stable employment with clear retention programs
- GenAI practice is a growth priority -- actual AI work, not just project labels
- Large team means mentorship opportunities and diverse project exposure

CONS / RED FLAGS

- IT services firm -- some roles become client-managed outsourcing rather than innovation
- Salary below pure-play product companies (Anthropic, Google DeepMind) by 30-50%
- Large matrix organisation -- career progression involves navigating multiple hierarchies
- Client-driven timelines mean variable hours during delivery crunch
- "Mid-level" label may limit initial compensation; negotiate to get senior band recognition
- Multi-geography calls may require early morning or late evening availability

RECENT NEWS & DEVELOPMENTS

- NTT DATA announced expansion of AI and GenAI practice in India (2025) -- Pune and Bangalore as key hubs
- Partnership with NVIDIA for AI infrastructure solutions
- Active investment in LLM deployment for enterprise clients across BFSI and healthcare
- NTT Group's \$3.6B annual R&D includes significant AI infrastructure investment
- Hiring aggressively for GenAI roles in India delivery centres through 2025-2026

TECH STACK (NTT DATA GenAI practice)

- LLMs: OpenAI GPT-4, Azure OpenAI, Anthropic Claude, open-source (LLaMA, Mistral)
- Frameworks: LangGraph, LangChain, CrewAI, Semantic Kernel, AutoGen
- Vector DBs: Pinecone, Qdrant, FAISS, Chroma
- Cloud: Azure (primary for enterprise clients), AWS, GCP
- Evaluation: LangSmith, RAGAS, custom eval frameworks

- Deployment: Docker, Kubernetes, Azure ML, AWS SageMaker
- Languages: Python (primary), occasional Java/TypeScript

NEGOTIATION INTEL

- NTT DATA has structured salary bands but 10-15% flex exists for experienced GenAI candidates
- Agentic AI + LangGraph + MCP is a scarce profile -- use this to push for senior band placement
- Target ask: ■28-34 LPA given 5+ years and agentic specialisation
- Negotiate joining bonus if base is capped -- typically ■1-2L available
- Ask specifically about: project allocation (avoid bench), GenAI practice team vs. generic delivery, certification budget
- Use the published checklist against them: "I tick every item on your resume selection criteria -- the salary should reflect a senior-band placement"
- Ask about the specific client or project: "Is this a new GenAI engagement or an existing account?" -- determines scope and growth trajectory

SECTION 3 — INTERVIEW PREPARATION KIT

TECHNICAL QUESTIONS (likely to be asked)

1. Walk me through a production agentic system you've built. What was the graph structure, how did the supervisor route, and what happened when an agent failed? -- Tests: Real LangGraph depth; answer with Contract Intelligence System architecture
2. What is hybrid search and how do you implement it alongside a vector store? -- Tests: Retrieval fundamentals; answer: dense vector search (semantic) + sparse BM25 (keyword) combined with weighted fusion or RRF; your cross-encoder reranking sits on top
3. You have PII redaction and guardrails on your CV. Walk me through what you built and how it works at runtime. -- Tests: Safety/production thinking; answer with your input/output guardrail filters and PII redaction in Enterprise Q&A Agent
4. How do you evaluate a GenAI system before showing results to a client? What metrics do you track? -- Tests: Evaluation maturity; answer: RAGAS (faithfulness, context precision, answer relevance), LLM-as-judge, hallucination rate, P95 latency -- you built this
5. What is the difference between LangGraph and CrewAI? When would you use each? -- Tests: Framework-level reasoning; LangGraph = stateful graph with conditional routing; CrewAI = role-based agents with sequential/hierarchical process
6. How do you handle agent memory -- both short-term within a conversation and long-term across sessions? -- Tests: Memory architecture; short-term = conversation buffer / sliding window; long-term = vector store with memory retrieval (semantic similarity on past interactions)
7. What chunking strategies have you used and how do you choose the right one? -- Tests: RAG engineering depth; answer: fixed-size, semantic, sentence-window, hierarchical -- choice depends on document type and query granularity
8. How do you implement feedback loops in a production RAG system to improve data quality over time? -- Tests: Production thinking; answer: collect user feedback signals, low-faithfulness queries flagged for review, re-ingestion pipeline on updated docs
9. Walk me through your LLM evaluation system. What dimensions did you measure and why those specifically? -- Tests: Evaluation design thinking; answer with your LLM Evaluator: accuracy (LLM-as-judge), hallucination, P95 latency, cost-per-query, instruction following
10. NTT DATA clients span multiple geographies. How have you communicated complex AI architecture to non-technical stakeholders? -- Tests: Consulting communication ability; answer with Optimum client delivery stories

BEHAVIOURAL QUESTIONS (NTT DATA-specific)

1. NTT DATA serves Fortune 100 clients -- tell me about a time you delivered under tight client deadlines with shifting requirements.

2. You'll work with stakeholders across US, Europe, and India. How do you manage multi-timezone communication on a delivery project?
3. Describe a time when a GenAI system you built had unexpected behaviour in production. How did you diagnose and fix it?
4. NTT DATA values both innovation and reliability. How do you balance experimenting with new frameworks vs. shipping stable solutions to clients?
5. Tell me about a time you had to push back on a client requirement because it wasn't technically feasible or would compromise quality.

STAR STORY SUGGESTIONS

Story 1: "Production agentic system with MCP"

- Situation: Needed automated, reliable legal contract analysis -- manual review was slow, inconsistent, and expensive
- Task: Design and build a multi-agent pipeline that could analyse contracts, identify risks, and benchmark against regulatory rules in real time
- Action: Built 6-agent LangGraph pipeline with supervisor routing and parallel fan-out; implemented 4 MCP stdio servers so agents could query clause libraries and regulatory rules mid-analysis; added feedback loop where risk agent is called again with corrective context when score is too high
- Result: System produces structured HTML reports with 0-100 risk score, missing clause detection, obligation mapping, compliance flags; works with both OpenAI and Claude
- Best used for: "Tell me about an agentic system you built", "Multi-agent architecture", "Complex project"

Story 2: "Production RAG with guardrails and evaluation"

- Situation: Enterprise clients needed a document Q&A system they could trust -- not just one that sometimes worked
- Task: Build a production RAG pipeline with proper evaluation, guardrails, and observability
- Action: Implemented 5-tier PDF extraction, cross-encoder reranking, HyDE, RAGAS evaluation framework, PII redaction, input/output guardrail filters, and L1-L4 quality indicators per query
- Result: System is measurable -- clients can see faithfulness score, relevance score, and hallucination detection per response; SSE streaming API with Docker Compose deployment
- Best used for: "RAG experience", "Guardrails", "Production AI thinking"

Story 3: "Client delivery across domains at Optimum"

- Situation: Multiple enterprise clients needed AI solutions across healthcare, BFSI, and industrial -- each with different requirements and technical levels
- Task: Deliver 3-5 successful POCs from requirements gathering through to production deployment
- Action: Managed full lifecycle for each -- scoped requirements with client, built models, evaluated against KPIs, deployed, and handed over with documentation
- Result: 95% accuracy on medical imaging, 92% on quality detection, working RAG chatbot -- all delivered and in use by clients

- Best used for: "Client delivery", "Multiple POCs", "Stakeholder communication"

SALARY NEGOTIATION STRATEGY

- Target ask: ■28-34 LPA
- Walk-away: ■22 LPA
- Script: "Given my LangGraph and MCP hands-on experience -- including a 6-agent production pipeline and 4 custom MCP servers -- alongside RAGAS-evaluated production RAG and 5+ years of Python, I'm targeting ■28-34 LPA. Your own resume checklist shows you're looking for exactly this profile, and I clear every criterion."
- Leverage: NTT DATA published their resume checklist -- use it. You tick every box, which justifies senior-band placement even at a mid-level role title.
- Ask explicitly: "What band does this role sit in and is there room to place at the top of that band given my Agentic AI specialisation?"
- Don't accept on the spot. Ask for 24 hours to review the written offer.

SMART QUESTIONS TO ASK

1. "Your resume checklist mentions 3-5 client POCs -- which specific client domain would this role serve first? BFSI, healthcare, or something else?"
2. "What Agentic AI frameworks is the team currently deploying in production -- LangGraph, CrewAI, or Semantic Kernel?"
3. "Is the GenAI practice a centralised Centre of Excellence or are engineers embedded in specific client accounts?"
4. "What does the multi-geography collaboration look like in practice -- which time zones are the primary client stakeholders in?"
5. "NTT DATA has a partnership with NVIDIA for AI infrastructure -- does this role interact with that work at all?"

KEY TALKING POINTS

1. "I clear every item on your published resume selection checklist: 2 delivered GenAI/Agentic solutions, 5+ years Python, 3-5 client POCs -- this isn't a borderline fit, it's a direct match"
2. "LangGraph, CrewAI, and MCP are all on my CV with real shipped code -- not tutorials, not prototypes, production systems"
3. "I evaluate my own systems with RAGAS and LLM-as-judge and build guardrails and PII redaction into production -- I care about reliability, not just demos"