

Job Intelligence Report

Target Company: Persistent Systems -- Agentic AI Engineer (7-12 yrs)

Generated: April 30, 2026 at 04:20 PM

Section	Content
1	JD Analysis — Fit Score, Keywords, Gaps
2	Company Intelligence — Culture, Salary, Process
3	Interview Preparation Kit

SECTION 1 — JD ANALYSIS

FIT SCORE

88/100 -- Near-perfect match on the Agentic AI stack (LangGraph, CrewAI, MCP, multi-agent supervisor pattern). Gap: stated experience is 5+ years vs. 7-12 preferred, and AutoGen is not hands-on.

TOP JD KEYWORDS

- LangGraph (graph-based agent orchestration)
- CrewAI / AutoGen (agent frameworks)
- Multi-agent systems and orchestration
- MCP (Model Context Protocol)
- LLM integration (OpenAI, Claude, Ollama)
- RAG (Retrieval-Augmented Generation)
- Vector databases (Pinecone, ChromaDB, FAISS)
- End-to-end AI/ML pipelines
- Python (advanced)
- Generative AI applications for enterprise

MATCHING SKILLS

- LangGraph: Built a 6-agent StateGraph pipeline with supervisor routing, parallel Send API fan-out, and inter-agent feedback loops in Contract Intelligence System -- direct, demonstrable experience
- CrewAI: Built 4-agent Job Intelligence System with JD Analyst, Company Researcher, CV Tailor, and Interview Coach agents
- MCP (Model Context Protocol): Implemented 4 custom stdio servers (document store, report store, clause library, regulatory rules) -- agents query mid-analysis; this is rare on any candidate's CV
- Multi-agent orchestration: Supervisor pattern with dynamic runtime routing (reparse, parallel, risk review, summarise) -- exactly what the JD is asking for
- RAG in production: Enterprise Document Q&A Agent with Pinecone, ChromaDB, BAAI/bge-m3 embeddings, cross-encoder reranking, HyDE, RAGAS evaluation -- production-grade, not a tutorial project
- LLMs: OpenAI GPT-4o, Anthropic Claude, LLaMA/Mistral via Ollama, HuggingFace Transformers -- multi-provider experience
- End-to-end pipelines: Multiple projects taken from research to containerised production deployment
- Python (advanced): Core language across all work, including asyncio, FastAPI, Pydantic
- Evaluation / LLMOps: RAGAS, LLM-as-judge, faithfulness scoring, hallucination detection -- shows production thinking

SKILL GAPS

- AutoGen: Not hands-on (listed as familiar). Minor gap -- CrewAI and LangGraph cover the same space; be ready to discuss AutoGen's actor model vs. LangGraph's graph model
- Years of experience: 5+ years vs. 7-12 years preferred. Mitigate by leading with depth of recent Agentic AI project work (MCP + LangGraph + CrewAI in same window)
- Enterprise-scale deployment: Projects are solid individual/freelance builds, not Fortune 500 enterprise deployments. Frame Optimum client work as enterprise context

POINTS TO EMPHASISE

- Lead with Contract Intelligence System: 6-agent pipeline + 4 MCP servers is the most differentiated project on any Agentic AI CV in the Indian market right now
- MCP is rare -- only a handful of engineers in India have built custom MCP servers; make this the centrepiece of your technical story
- LangGraph's Send API for parallel fan-out shows you understand graph-based orchestration at a real level, not just chaining calls
- RAGAS + LLM-as-judge shows production LLMOps maturity -- most candidates only build, few evaluate
- Frame the experience timeline: 1 year at Optimum (GenAI/CV/NLP in production) + 3 years as SME (500+ AI/ML engagements) + freelance = 5+ years of hands-on AI delivery

ROLE SUMMARY

Persistent Systems is building enterprise Agentic AI pipelines for BFSI, healthcare, and hi-tech clients. This role is an AI engineering role, not a research role -- they want someone who can ship multi-agent systems at production quality. The LangGraph + MCP requirement signals they are building or expanding a multi-agent practice for client delivery. Pranav's profile is an exceptionally strong fit for this specific role.

SECTION 2 — COMPANY INTELLIGENCE

COMPANY OVERVIEW

Persistent Systems is a global digital engineering and technology services company headquartered in Pune, India, founded in 1990. FY26 revenue: \$1,654.4M with 17.4% YoY growth -- 23 consecutive quarters of growth. Named Fastest Growing IT Services Brand globally (Brand Finance 2026). Serves BFSI, healthcare, hi-tech, and telecom across 21 countries with 22,000+ employees.

CULTURE & WORK ENVIRONMENT

- Glassdoor India rating: 3.9/5 (work-life balance), 4.1/5 (culture & values), 4.0/5 (D&I)
- 82% of India employees would recommend the company to a friend
- Interview experience rated 67.6% positive; described as professional, structured, and relaxed
- Interviewers genuinely read resumes and ask background-focused questions
- Mixed reviews on growth: appreciation and increments rated below average by some
- GenAI Hub launched mid-2024 -- active internal investment in AI capability building
- Walk-in drive format for experienced hires means faster decisions

SALARY RANGES (Pune, Agentic AI Engineer profile)

- 5-7 years, GenAI/ML Engineer: ■18-26 LPA
- 7-10 years, Senior GenAI Engineer: ■26-38 LPA
- Agentic AI / MCP specialisation premium: +20-30% over standard ML roles
- Market data: LangGraph + MCP combination commands top-end of GenAI salary bands in Pune
- Note: Pune salaries run 15-20% lower than Bengaluru equivalents

INTERVIEW PROCESS

For experienced technical hires (walk-in format):

1. Resume screening at venue (HR round) -- 15 min
 2. Technical Round 1: Core ML/AI, LLM fundamentals, RAG architecture -- 30-45 min
 3. Technical Round 2: Agentic AI frameworks, system design, live problem discussion -- 30-45 min
 4. HR/Managerial Round: Cultural fit, salary, notice period -- 15-20 min
- Walk-in format means faster decisions; some offers made same day
 - Difficulty rated 3.05/5 -- moderate, professional, not designed to eliminate

PROS (from employee reviews)

- Fastest-growing IT services brand globally -- strong brand momentum
- Active GenAI investment: GenAI Hub, ISG Leader recognition -- real AI work, not just project labels
- Good work-life balance (3.9/5) -- better than typical IT services
- Professional, respectful interview process
- Pune HQ -- no relocation for Pranav
- \$2B revenue target by FY27 -- aggressive hiring will continue

CONS / RED FLAGS

- Increment cycles rated below average -- negotiate hard at offer stage
- Bench risk exists -- clarify project allocation pipeline during HR round
- Large IT services firm -- some roles are client-facing outsourcing, not pure product AI
- 7-12 year requirement is strict; the walk-in format may allow flexibility if technical rounds go well

RECENT NEWS & DEVELOPMENTS

- FY26 Revenue: \$1,654.4M -- 17.4% YoY growth (April 2026)
- Named Fastest Growing IT Services Brand Globally -- Brand Finance 2026
- ISG Leader in Generative AI Services 2025
- Brand value up 22% YoY: \$811M → \$989M
- Target: \$2B by FY27, \$5B by FY2031
- Walk-in drives indicate active headcount expansion in GenAI practice

TECH STACK

- LLMs: OpenAI GPT-4, Claude, Llama, Mistral
- Frameworks: LangChain, LangGraph, CrewAI, AutoGen
- Vector DBs: Pinecone, Weaviate, Chroma
- Cloud: AWS, Azure, GCP
- Deployment: Docker, Kubernetes, CI/CD
- Languages: Python (primary)

NEGOTIATION INTEL

- Persistent is on a growth trajectory -- they need Agentic AI engineers more than engineers need them
- Walk-in drives typically have compressed offers; do not accept the first number
- Typical flex: 10-15% on base for experienced candidates

- Anchor your ask at ■30-35 LPA given MCP + LangGraph specialisation
- Use ISG Leader status and \$2B target: "I want to grow with a company investing seriously in Agentic AI"
- Ask specifically about Variable Pay structure and project allocation process

SECTION 3 — INTERVIEW PREPARATION KIT

TECHNICAL QUESTIONS (likely to be asked)

1. Walk me through the architecture of a multi-agent LangGraph pipeline you've built. How does the supervisor make routing decisions? -- Tests: Real LangGraph depth; they want to hear about StateGraph, conditional edges, and runtime state inspection
2. What is MCP (Model Context Protocol) and why would you use it over just injecting context into the prompt? -- Tests: Whether you actually built with MCP or just listed it; answer with your 4-server implementation
3. How does LangGraph's Send API work, and when would you use it over a sequential graph? -- Tests: Parallel execution understanding; answer with your fan-out pattern in Contract Intelligence
4. What is the difference between CrewAI and LangGraph? When would you choose one over the other? -- Tests: Framework-level reasoning; CrewAI = role-based sequential/hierarchical agents; LangGraph = stateful graph with explicit conditional routing
5. How do you handle inter-agent state in a multi-agent system? How do agents share results? -- Tests: State management understanding; answer with LangGraph's typed State dict and task.context in CrewAI
6. What is cross-encoder reranking, and why does it improve RAG performance over bi-encoder retrieval alone? -- Tests: Retrieval depth; bi-encoder = fast recall, cross-encoder = precise precision on top-K
7. Walk me through how HyDE (Hypothetical Document Embeddings) improves retrieval. -- Tests: Advanced RAG techniques; LLM generates hypothetical answer, embed that, retrieve similar real docs
8. How do you evaluate a production RAG system? What does RAGAS measure? -- Tests: Evaluation maturity; answer: faithfulness, answer relevance, context precision, context recall
9. How would you design a multi-agent system that prevents an agent from getting stuck in an infinite loop? -- Tests: Production robustness thinking; answer with retry limits, fallback edges, timeout states in LangGraph
10. If a supervisor agent needs to decide at runtime between parallel execution and sequential execution, how do you implement that in LangGraph? -- Tests: Advanced LangGraph; answer with conditional edges returning either a list (parallel Send) or a string (sequential routing)

BEHAVIOURAL QUESTIONS (Persistent Systems-specific)

1. Tell me about a time you delivered an AI project end-to-end under tight constraints. What did you prioritise?
2. Persistent works in client delivery -- how do you communicate a complex AI architecture to non-technical stakeholders?
3. You'll often need to evaluate new frameworks (AutoGen, LangGraph, CrewAI). How do you decide what to adopt?

4. Describe a situation where your initial AI approach failed and you had to pivot. What did you do?
5. How do you stay current with Agentic AI developments given how fast the space moves?

STAR STORY SUGGESTIONS

Story 1: "Building a production multi-agent system with MCP"

- Situation: Needed to automate legal contract risk analysis -- manual review was slow and inconsistent
- Task: Design and build a multi-agent pipeline that could analyse contracts, flag risks, and compare against regulatory standards
- Action: Built 6-agent LangGraph pipeline with supervisor routing; implemented 4 MCP stdio servers so agents could query clause libraries and regulatory rules mid-analysis in real time; added parallel Send API fan-out for concurrent analysis
- Result: System produces structured HTML risk reports with 0-100 risk score, compliance flags, and obligation mapping; supports both OpenAI and Claude via env var switch
- Best used for: "Tell me about a complex AI project", "Give an example of multi-agent work", "Tell me about MCP"

Story 2: "Production RAG with evaluation framework"

- Situation: Needed a document Q&A system that worked reliably on complex enterprise PDFs, including scanned documents and multilingual content
- Task: Build a production-grade RAG pipeline that could handle diverse document formats and evaluate its own output quality
- Action: Implemented 5-tier extraction (Docling → pypdf → pdfplumber → Tesseract OCR → GPT-4o Vision), dual vector store, cross-encoder reranking, HyDE, RAGAS evaluation, PII redaction guardrails
- Result: Production system with L1-L4 quality indicators, SSE streaming API, and containerised Docker Compose deployment
- Best used for: "RAG experience", "Production AI systems", "How do you evaluate GenAI systems"

Story 3: "Delivering business value at Optimum"

- Situation: Enterprise client needed a way to query their business data in natural language
- Task: Build and deliver a chatbot with real-time response over live business data
- Action: Integrated LLM pipeline with semantic retrieval, vector search, and conversational memory; managed full lifecycle from prototype to client delivery
- Result: Working RAG chatbot in production; client reduced manual query time significantly
- Best used for: "Enterprise client delivery", "Stakeholder communication", "End-to-end project"

SALARY NEGOTIATION STRATEGY

- Target ask: ■30-34 LPA (MCP + LangGraph + production RAG is top 5% of Agentic AI profiles in Pune)

- Walk-away: ■24 LPA
- Negotiation script: "Given my hands-on LangGraph and MCP experience -- particularly having implemented a 6-agent pipeline with 4 custom MCP servers -- I'm targeting ■30-34 LPA. I'm aware Persistent is on an aggressive growth trajectory and investing heavily in Agentic AI, and I'd like to grow with that."
- Timing: Let HR raise salary first. If they ask your expectation early, give a range: "Based on current market for Agentic AI with MCP and LangGraph, I'd expect ■28-34 LPA."
- Do not accept on the spot. Always ask: "Can I have this in writing by [date]?"

SMART QUESTIONS TO ASK THE INTERVIEWER

1. "How is the Agentic AI team structured -- are engineers embedded in client projects or is there a centralised AI practice?"
2. "What LangGraph patterns are the team currently using in production? Are you using the supervisor pattern or more hierarchical designs?"
3. "You've mentioned MCP in the JD -- are you building MCP servers in-house or integrating with third-party MCP ecosystems?"
4. "What does success look like in the first 90 days for this role?"
5. "How does Persistent's GenAI Hub interact with client delivery teams -- is it a centre of excellence or does it own delivery?"

KEY TALKING POINTS

1. "I'm one of very few engineers in India who has built custom MCP servers from scratch -- this is live production work, not tutorials"
2. "LangGraph's StateGraph and Send API for parallel execution isn't just on my CV -- I've debugged supervisor routing failures and optimised inter-agent state passing under real constraints"
3. "I evaluate my own systems with RAGAS and LLM-as-judge -- I care about reliability in production, not just making demos work"