

Job Intelligence Report

Target Company: Persistent Systems -- GenAI Engineer (5-8 yrs)

Generated: April 30, 2026 at 04:20 PM

Section	Content
1	JD Analysis — Fit Score, Keywords, Gaps
2	Company Intelligence — Culture, Salary, Process
3	Interview Preparation Kit

SECTION 1 — JD ANALYSIS

FIT SCORE

92/100 -- Exceptional match. LangGraph, MCP, multi-agent, RAG, LLMs all explicitly present and demonstrated with real shipped projects. Experience range (5+ years) fits squarely. The small gap is stating ~3 years of dedicated GenAI experience vs. the preferred "3 years" -- arguable depending on how you count.

TOP JD KEYWORDS

- LangGraph (explicitly preferred framework -- #1 priority)
- MCP (Model Context Protocol -- listed as must-have)
- Multi-agent systems and orchestration
- Agentic AI (6+ months hands-on required)
- Generative AI (~3 years preferred)
- AWS (working understanding)
- Full AI/ML lifecycle project management
- Python (advanced)
- Core ML + Deep Learning foundations
- RAG and vector databases

MATCHING SKILLS

- LangGraph: Contract Intelligence System built on LangGraph StateGraph with supervisor routing, parallel Send API, and conditional edges -- direct, demonstrable, recent experience
- MCP: 4 custom stdio servers implemented from scratch for Contract Intelligence -- this is what the JD calls out as must-have, and it is rare
- Multi-agent systems: Supervisor pattern with dynamic runtime routing, parallel fan-out, inter-agent feedback loops; also 4-agent CrewAI pipeline in Job Intelligence System
- Agentic AI (6+ months): Contract Intelligence + Job Intelligence System together = substantial hands-on Agentic AI work
- GenAI (~3 years): 1 year at Optimum (RAG, LLMs, production GenAI), 1 year of Agentic AI project development, freelance financial AI -- totals ~2-2.5 years of direct GenAI work; frame honestly as "growing rapidly in the space"
- RAG (production): Enterprise Document Q&A with Pinecone, ChromaDB, HyDE, cross-encoder reranking, RAGAS
- AWS: EC2, S3 working knowledge -- meets "working understanding" bar explicitly
- Full lifecycle project management: Research → prototype → evaluation → client delivery across Optimum projects
- Python (advanced): Core language; asyncio, FastAPI, Pydantic, Plotly all demonstrated

- Core ML/DL: PyTorch, TensorFlow, Keras, HuggingFace -- solid foundational stack

SKILL GAPS

- ~3 years dedicated GenAI experience: Direct GenAI work is ~2-2.5 years; frame as "concentrated recent growth in Agentic AI" -- the depth of Contract Intelligence + Enterprise RAG compensates for timeline
- Project management language: JD emphasises stakeholder communication and sprint management; add explicit PM framing in interview answers
- Scale: Projects are strong individual builds; Persistent will want to see evidence of working in larger team/enterprise contexts -- use Optimum client delivery stories

POINTS TO EMPHASISE

- Lead immediately with MCP: "I've built 4 custom MCP studio servers" -- this is the single rarest skill on the JD and in the Indian market
- LangGraph is the JD's #1 preferred framework and you have a complete 6-agent pipeline built on it -- this is the strongest possible answer to their top requirement
- The combination of LangGraph + MCP + CrewAI shows breadth across the Agentic AI stack, not just knowledge of one tool
- RAGAS + LLM-as-judge in production shows evaluation maturity most candidates lack
- GenAI experience: 1 year production (Optimum) + concentrated project development = strong recency signal even if total calendar years are 2+

ROLE SUMMARY

This GenAI Engineer role (5-8 years) is Persistent's front-line AI engineering position for their GenAI practice. Unlike a research role, it requires shipping -- they want someone who can manage a GenAI project from brief through to client delivery, using LangGraph and MCP as core tools. The explicit call-out of MCP in a job description is still unusual in the Indian market, which means Persistent is specifically building an MCP-enabled agentic capability. Pranav's profile is arguably the best possible match for this specific requirement combination.

SECTION 2 — COMPANY INTELLIGENCE

COMPANY OVERVIEW

Persistent Systems is a global digital engineering and technology services company headquartered in Pune, India, founded in 1990. FY26 revenue: \$1,654.4M with 17.4% YoY growth -- 23 consecutive quarters of growth. Named Fastest Growing IT Services Brand globally (Brand Finance 2026). Serves BFSI, healthcare, hi-tech, and telecom across 21 countries with 22,000+ employees.

CULTURE & WORK ENVIRONMENT

- Glassdoor India rating: 3.9/5 (work-life balance), 4.1/5 (culture & values), 4.0/5 (D&I)
- 82% of India employees would recommend the company to a friend
- Interview experience rated 67.6% positive; described as professional, structured, and relaxed
- Interviewers genuinely read resumes and ask background-focused questions
- Mixed reviews on growth: appreciation and increments rated below average by some
- GenAI Hub launched mid-2024 -- active internal investment in AI capability building
- Walk-in drive format for experienced hires means faster decisions

SALARY RANGES (Pune, Agentic AI Engineer profile)

- 5-7 years, GenAI/ML Engineer: ■18-26 LPA
- 7-10 years, Senior GenAI Engineer: ■26-38 LPA
- Agentic AI / MCP specialisation premium: +20-30% over standard ML roles
- Market data: LangGraph + MCP combination commands top-end of GenAI salary bands in Pune
- Note: Pune salaries run 15-20% lower than Bengaluru equivalents

INTERVIEW PROCESS

For experienced technical hires (walk-in format):

1. Resume screening at venue (HR round) -- 15 min
 2. Technical Round 1: Core ML/AI, LLM fundamentals, RAG architecture -- 30-45 min
 3. Technical Round 2: Agentic AI frameworks, system design, live problem discussion -- 30-45 min
 4. HR/Managerial Round: Cultural fit, salary, notice period -- 15-20 min
- Walk-in format means faster decisions; some offers made same day
 - Difficulty rated 3.05/5 -- moderate, professional, not designed to eliminate

PROS (from employee reviews)

- Fastest-growing IT services brand globally -- strong brand momentum
- Active GenAI investment: GenAI Hub, ISG Leader recognition -- real AI work, not just project labels
- Good work-life balance (3.9/5) -- better than typical IT services
- Professional, respectful interview process
- Pune HQ -- no relocation for Pranav
- \$2B revenue target by FY27 -- aggressive hiring will continue

CONS / RED FLAGS

- Increment cycles rated below average -- negotiate hard at offer stage
- Bench risk exists -- clarify project allocation pipeline during HR round
- Large IT services firm -- some roles are client-facing outsourcing, not pure product AI
- 7-12 year requirement is strict; the walk-in format may allow flexibility if technical rounds go well

RECENT NEWS & DEVELOPMENTS

- FY26 Revenue: \$1,654.4M -- 17.4% YoY growth (April 2026)
- Named Fastest Growing IT Services Brand Globally -- Brand Finance 2026
- ISG Leader in Generative AI Services 2025
- Brand value up 22% YoY: \$811M → \$989M
- Target: \$2B by FY27, \$5B by FY2031
- Walk-in drives indicate active headcount expansion in GenAI practice

TECH STACK

- LLMs: OpenAI GPT-4, Claude, Llama, Mistral
- Frameworks: LangChain, LangGraph, CrewAI, AutoGen
- Vector DBs: Pinecone, Weaviate, Chroma
- Cloud: AWS, Azure, GCP
- Deployment: Docker, Kubernetes, CI/CD
- Languages: Python (primary)

NEGOTIATION INTEL

- Persistent is on a growth trajectory -- they need Agentic AI engineers more than engineers need them
- Walk-in drives typically have compressed offers; do not accept the first number
- Typical flex: 10-15% on base for experienced candidates

- Anchor your ask at ■30-35 LPA given MCP + LangGraph specialisation
- Use ISG Leader status and \$2B target: "I want to grow with a company investing seriously in Agentic AI"
- Ask specifically about Variable Pay structure and project allocation process

SECTION 3 — INTERVIEW PREPARATION KIT

TECHNICAL QUESTIONS (likely to be asked)

1. You have LangGraph and MCP on your CV -- can you describe exactly what MCP is and why you chose to use it in your project? -- Tests: Whether MCP is real experience; answer with your 4 studio servers and what each one does
2. In LangGraph, what is the difference between a node and an edge? How do you implement conditional routing? -- Tests: LangGraph fundamentals; answer with StateGraph, conditional_edges, and your supervisor routing function
3. How do you manage shared state across 6 agents in a LangGraph pipeline? What does your State object look like? -- Tests: State management depth; answer with typed TypedDict State, channels, and how agents read/write to it
4. What is HyDE and how does it improve retrieval quality? When would you NOT use it? -- Tests: Advanced RAG; HyDE adds latency -- don't use for short factual queries where hypothetical answer diverges
5. Walk me through the 5-tier PDF extraction pipeline you built. Why 5 tiers? -- Tests: Production engineering thinking; each tier handles a failure mode of the previous
6. How do you detect and reduce hallucination in a production RAG system? -- Tests: Faithfulness scoring, LLM-as-judge, context grounding -- answer with RAGAS faithfulness metric
7. What is cross-encoder reranking and why is it better than just using bi-encoder retrieval? -- Tests: Retrieval fundamentals; bi-encoder = recall, cross-encoder = precision
8. How would you manage a GenAI project sprint with a client -- what does your delivery lifecycle look like? -- Tests: PM ability; answer with scoping → prototype → eval → iteration → deployment flow from Optimum experience
9. If a client asks you to add a new agent to an existing LangGraph pipeline, what are the risks and how do you manage them? -- Tests: Production thinking; state schema changes, edge re-wiring, regression on existing paths
10. How do you choose between Pinecone and ChromaDB for a new project? -- Tests: Vector DB practical knowledge; answer: Pinecone for cloud/production scale, ChromaDB for local dev -- you built dual-store supporting both

BEHAVIOURAL QUESTIONS (Persistent Systems GenAI practice)

1. This role involves both building AI systems AND communicating progress to clients. Tell me about a time you translated complex AI work into business language.
2. LangGraph and MCP are both very new. How do you evaluate a new framework before adopting it on a client project?

3. Describe a GenAI project where the first approach didn't work and how you recovered.
4. How do you stay on top of the Agentic AI space -- what are you reading or building right now?
5. Persistent invests in employee certifications. What would you want to learn or certify in next?

STAR STORY SUGGESTIONS

Story 1: "MCP server implementation"

- Situation: Needed agents to access external knowledge bases (clause library, regulatory rules) mid-analysis without re-ingesting everything into context
- Task: Design a protocol-based access layer that agents could query dynamically
- Action: Built 4 MCP stdio servers, each with defined tool schemas; agents call these like function tools but the data lives outside the model context
- Result: Agents retrieve only what they need when they need it -- no context bloat, clean separation of concerns
- Best used for: MCP questions, system design questions, "tell me about your most innovative project"

Story 2: "RAG evaluation pipeline"

- Situation: After building the Enterprise Document Q&A agent, needed to prove it was actually reliable before presenting to stakeholders
- Task: Build an evaluation framework that measures faithfulness, relevance, and hallucination rate
- Action: Integrated RAGAS framework, implemented LLM-as-judge for faithfulness scoring, built L1-L4 quality indicator system in the Streamlit UI, added PII redaction guardrails
- Result: System has measurable quality metrics per query -- can demonstrate to clients exactly how reliable it is
- Best used for: LLMOps, evaluation, production thinking questions

Story 3: "Full lifecycle delivery at Optimum"

- Situation: Client needed a business query chatbot over their own data -- no existing AI infrastructure
- Task: Build and deliver a production RAG chatbot from scratch, managing the full lifecycle
- Action: Scoped requirements, built RAG pipeline with semantic retrieval and conversational memory, containerised with Docker, delivered to client with documentation
- Result: Client has working production chatbot; I managed from initial briefing through deployment and handover
- Best used for: PM questions, client delivery questions, end-to-end lifecycle

SALARY NEGOTIATION STRATEGY

- Target ask: ■28-32 LPA (strong match, 5-8 yr band, MCP specialisation)
- Walk-away: ■22 LPA

- Script: "Given my hands-on LangGraph and MCP experience -- which are the two specific frameworks called out in your JD -- I'm targeting ■28-32 LPA. I've seen current market rates and this aligns with Agentic AI profiles in Pune at this experience level."
- Timing: Wait for HR to ask. If they press for a number, give the range.
- Always ask about the full package: fixed, variable, joining bonus, and cert reimbursement

SMART QUESTIONS TO ASK

1. "The JD calls out MCP specifically -- are you building a proprietary MCP server ecosystem or integrating with open MCP registries?"
2. "How do you structure GenAI project sprints -- what does a typical 2-week cycle look like for this team?"
3. "What is the split between building new pipelines vs. maintaining and improving existing client deployments?"
4. "How does the GenAI Hub team collaborate with this role -- are they a resource pool or do they own delivery?"
5. "What does the onboarding process look like -- is there a project I'd be assigned to in week one?"

KEY TALKING POINTS

1. "LangGraph and MCP are your two top requirements -- I've shipped production systems using both, including 4 custom MCP stdio servers that agents query dynamically"
2. "I evaluate my AI systems with RAGAS and LLM-as-judge -- I care about production reliability, not just getting a demo to work"
3. "I've managed full AI project lifecycles with real stakeholders -- from business briefing through containerised production delivery"